

A decorative banner featuring five large, bold, blue letters ('I', 'N', 'D', 'E', 'X') arranged horizontally. Each letter is centered within a white square frame with a thin blue border. The letters are slightly tilted, giving them a dynamic appearance.

NAME: Stuti Uniyal STD.: _____ SEC.: D ROLL NO.: _____ SUB.: Machine Learning

WEEK - 1
PANDAS (Reading CSV)

import pandas as pd

df2 = pd.read_csv("/content/sample-data/california_housing-test.csv")

df2.head()

OUTPUT:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
0	-122.05	37.37	27.0	3885.0	661.0
	population	households	median_income	median_house_value	
0	1537.0	606.0	6.0085	344700.0	

df2.to_csv("hi.csv")

col_names = ['sepal-length-in-cm', 'sepal-width-in-cm', 'petal-length-in-cm', 'petal-width-in-cm', 'class']

wel = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"

df = pd.read_csv(wel, names=col_names)

df.head()

OUTPUT:

	sepal_length-in-cm	sepal_width-in-cm	petal_length-in-cm
0	5.1	3.5	1.4

petal_width-in-cm class
0.2 Iris-setosa

df[0]

1. Performance Measure

2. Get the Data

2.1 Download the data

```
import os
import tarfile
import urllib
```

```
DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml2/master/"
HOUSING_PATH = os.path.join("data", "y1")
HOUSING_URL = DOWNLOAD_ROOT + "dataset/housing/housing.tgz"
```

```
def fetch_housing_data(housing_url=HOUSING_URL,
                      housing_path=HOUSING_PATH):
    os.makedirs(name=housing_path, exist_ok=True)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(url=housing_url, filename=tgz_path)
    housing_tgz = tarfile.open(name=tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()
```

```
fetch_housing_data()
```

```
import pandas as pd
```

```
def load_housing_data(housing_path=HOUSING_PATH):
    data_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(data_path)
```

2.2 Take a Quick Look at the Data Structure

```
housing = load_housing_data()
housing.head()
housing.info()
```

1. housing['ocean_proximity'].value_counts()
2. housing.describe()
2.3 PLOTTING GRAPHS
2. import matplotlib.pyplot as plt
3. import seaborn as sns
4. housing.hist(bins=50, figsize=(20, 15))
5. plt.show()

8.4 Create a Test set

```
import numpy as np
def split_train_test(data, test_ratio=0.2):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

train_set, test_set = split_train_test(data=housing)

```
len(train_set), len(test_set)
```

from zile import or32

```
def test_set_check(identifier, test_ratio=0.2):
    total_size = 2**32
    hex_repr = hex(int64(identifier))
    int_hex = hex(int(hex_repr[-2:], 16) < (test_ratio * total_size))
```

[test_set_check(i) for i in range(10)]

2.5 Co-ordinates to Index

```
def from_z_to_N(z):
    if z >= 0:
        n = 2**z
```

else :

$$n = -2 * 2 - 1$$

return n

def cantor_pairing(n1, n2) :

$$n = ((n_1 + n_2) * (n_1 + n_2 + 1)) / 2 + n_2$$

return n

def lat_lon_to_index(lat, lon) :

$$\text{lat, lon} = \text{int}(\text{lat} * 100), \text{int}(\text{lon} * 100)$$

$$\text{lat, lon} = \text{from_z_to_N}(\text{lat}), \text{from_z_to_N}(\text{lon})$$

$$\text{index} = \text{cantor_pairing}(\text{lat}, \text{lon})$$

return $n * \text{int}64 / \text{index}$

housing['id'] = housing.apply(lambda row: lat_lon_to_index(row['latitude'], row['longitude']), axis=1)

housing['id'].value_counts()

Day 28/3/24

Discover & Visualize the Data to Your Insights

```
strat_train_set.shape, strat_test_set.shape  
((16512, 10), (4128, 10))
```

```
strat_test_set.reset_index(), tofeather(frame='data')  
housing = strat_train_set.copy();  
housing.shape
```

Looking for correlations

```
corr_matrix = housing.corr()
```

```
corr_matrix['median_house_value'].sort_values(ascending=False)  
from pandas import plotting import scatter_matrix
```

~~attributes = ['median_house_value', 'median_income', 'total_rooms',
'housing_median_age']~~

```
scatter_matrix(frame=housing[attributes], figsize=(12, 8))  
plt.show()
```

WEEK 3 (PART-B)

Regressor model

```
regressor = LinearRegression()  
regressor.fit(x-train, y-train)
```

prediction result

```
y-pred-test = regressor.predict(x-test)
```

```
y-pred-train = regressor.predict(x-train)
```

prediction on training set

```
plt.scatter(x-train, y-train, color='lightcoral')
```

```
plt.plot(x-train, y-pred-train, color='firebrick')
```

```
plt.title('Salary vs Experience (Training Set),')
```

```
plt.xlabel('Years of Experience')
```

```
plt.ylabel('Salary')
```

```
plt.legend(['x-train/Pred(y-test)', 'x-train/y-train'])
```

```
title = 'Sal/Exp', loc='best', facecolor='white')
```

```
plt.box(False)
```

```
plt.show()
```

prediction on test set

```
plt.scatter(x-test, y-test, color='lightcoral')
```

```
plt.plot(x-train, y-pred-train, color='firebrick')
```

```
plt.title('Salary vs Experience (Test set),')
```

~~plt.xlabel('Years of Experience')~~~~plt.ylabel('Salary')~~~~plt.legend(['x-train/Pred(y-test)', 'x-train/y-train'])~~~~title = 'Sal/Exp', loc='best', facecolor='white')~~~~plt.box(False)~~~~plt.show()~~

Regressor coefficients and intercepts

```
print(f'Coefficient: {regressor.coef_3'})
```

```
print(f'Intercept: {regressor.intercept_3'})
```

OUTPUT

Coefficient: $[9318.57512673]$

Intercept: $[86780.09915063]$

~~Ansulay~~



WEEK - 4

```

import pandas
import matplotlib
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target
from sklearn.model_selection import train_test_split
x-train, x-test, y-train, y-test = train_test_split(
    X, y, test_size=0.2)

```

~~#~~

```

from sklearn.tree import DecisionTreeClassifier
clf = DecisionTreeClassifier(max_depth=3, min_samples_leaf=10,
                             random_state=1)
clf.fit(x-train, y-train)
y-pred = clf.predict(x-test)

```

```

from sklearn.metrics import accuracy_score
print(accuracy_score(y-test, y-pred))

```

```

from matplotlib.pyplot import rcParams
rcParams['figure.figsize'] = 80, 50

```

```

from sklearn.tree import plot_tree
plot_tree(clf)

```

OUTPUT

0.9666

$X[3] \leq 0.8$
 gini = 0.666
 samples = 100
 value = [43, 38, 39]

gini = 0.0
 samples = 43
 samples = 0
 value = [43, 0, 0]

$X[2] \leq 4.75$
 gini = 0.5
 samples = 77
 values = [0, 38, 39]

18/4/24

LAB-5

```
import pandas as pd
```

```
from matplotlib import pyplot as plt
```

```
%matplotlib inline
```

```
df = pd.read_csv ("content/drive/myDrive/insurance.csv")
```

```
df.head()
```

```
plt.scatter(df.age, df.bought-insurance, marker='+',  
           color='red')
```

```
from sklearn.model_selection import train-test-split
```

```
x-train, x-test, y-train, y-test = train-test-split(  
    df[['age']], df.bought-insurance,  
    train-size=0.7)
```

```
from sklearn.linear-model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(x-train, y-train)
```

```
y-predict = model.predict(x-test)
```

```
print(model.predict_proba(x-test))
```

```
print(model.predict(x-test))
```

```
model.score(x-test, y-test)
```

output:

```
[0.21 0.11 0.1]
```

```
0.777
```

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(x-train, y-train)  
print(model.predict(x-test))  
model.score(x-test, y-test)
```

OUTPUT:

0.58433

```
import math
```

```
def sigmoid(x):
```

```
    return 1/(1+math.exp(-x))
```

```
def prediction_function(age):
```

$z = 0.042 * \text{age} - 1.53$

$y = \text{sigmoid}(z)$

```
    return y
```

age = 35

```
print(prediction_function(age))
```

age = 43

~~print(prediction_function(age))~~

OUTPUT:

0.485004

0.56956

25/11/24

LAB-6

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.svm import train_test_split  
from sklearn.model_selection import confusion_matrix  
from sklearn.metrics import accuracy_score, classification_report  
from sklearn.neighbors import KNeighborsClassifier
```

```
colnames = ["sepal-length-in-cm", "sepal-width-in-cm",  
            "petal-length-in-cm", "petal-width-in-cm", "class"]  
dataset = pd.read_csv("https://archive.ics.uci.edu/ml/  
                         machine-learning-databases/iris/iris.data",  
                         header=None, names=colnames)
```

```
dataset.head()
```

```
dataset = dataset.replace({"class": {"Iris-setosa": 1,  
                                     "Iris-versicolor": 2, "Iris-virginica": 3}})
```

```
dataset.head()
```

```
iris = sns.load_dataset('iris')  
sns.set_style("whitegrid")  
sns.FacetGrid(iris, hue="species", height=6).  
map(plt.scatter, 'sepal-length', 'petal-length').  
add_legend()
```

```
X = dataset.iloc[:, :-1]
```

```
y = dataset.iloc[:, -1].values
```

```
x_train, x_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.25, random_state=0)
```

classifier = SVC(kernel='linear', random_state=0)
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)

$\begin{bmatrix} 13 & 0 & 0 \\ 0 & 15 & 1 \\ 0 & 0 & 9 \end{bmatrix}$

y_pred = classifier.predict(x_test)
print(y_pred)

print(classification_report(y_test, y_pred, target_names=['1', '2', '3']))

accuracy 0.97

classifier_2 = KNeighborsClassifier(n_neighbors=2)
classifier_2.fit(x_train, y_train)
y_pred = classifier_2.predict(x_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)

$\begin{bmatrix} 13 & 0 & 0 \\ 0 & 15 & 1 \\ 0 & 0 & 9 \end{bmatrix}$

y_pred = classifier_2.predict(x_test)

print(classification_report(y_test, y_pred, target_names=['1', '2', '3']))

accuracy 0.7
Date 9/5/24

LAB 8

Implementation of ANN using Back propagation for given values

```
import numpy as np
x = np.array([[0, 1], [1, 0], [1, 1], [0, 0]], dtype=float)
y = np.array([0, 1, 1, 0], dtype=float)
x = x / np.linalg.norm(x, axis=0)
y = y / 100
```

epoch = 5000

lr = 0.1

inputlayer_neurons = 2

hiddenlayer_neurons = 3

output_neurons = 1

wh = np.random.uniform(size=(inputlayer_neurons, hiddenlayer_neurons))

bh = np.random.uniform(size=(1, hiddenlayer_neurons))

wout = np.random.uniform(size=(hiddenlayer_neurons, output_neurons))

bout = np.random.uniform(size=(1, output_neurons))

def sigmoid(x):

return 1 / (1 + np.exp(-x))

def derivatives_sigmoid(x):

return x * (1 - x)

for i in range(epoch):

hinp1 = np.dot(x, wh)

hinp = hinp1 + bh

layer1_act = sigmoid(hinp)

output1_outinp1 = np.dot(layer1_act, wout)

outinp = outinp1 + bout

output = sigmoid(outinp)

$E_0 = y - \text{output}$

$\text{outgrad} = \text{derivatives - sigmoid (output)}$

$d\text{-output} = E_0 * \text{outgrad}$

$E_H = d\text{-output} \cdot \text{dot}(w_{\text{out}} \cdot T)$

$\text{hiddergrad} = \text{derivatives - sigmoid (hlayer_act)}$

$d\text{-hidden-layer} = E_H * \text{hiddergrad}$

$w_{\text{out}} + = \text{hlayer_act} \cdot T \cdot \text{dot}(d\text{-output}) * \text{lr}$

$v_{\text{h}} = X \cdot T \cdot \text{dot}(d\text{-hidden_layer}) * \text{lr}$

`print("Input : " + str(x))`

`print("Actual output : " + str(y))`

`print("Predicted output :\n", output)`

OUTPUT

Input:

$\begin{bmatrix} 0.6667 & 1. \\ 0.3334 & 0.556 \\ 1. & 0.6667 \end{bmatrix}$

Actual output :

$\begin{bmatrix} 0.92 \\ 0.26 \\ 0.89 \end{bmatrix}$

~~$\begin{bmatrix} 0.92 \\ 0.26 \\ 0.89 \end{bmatrix}$~~

Predicted output:

$\begin{bmatrix} 0.935 \\ 0.923 \\ 0.9339 \end{bmatrix}$

$\begin{bmatrix} 0.923 \\ 0.9339 \end{bmatrix}$

$\begin{bmatrix} 0.9339 \end{bmatrix}$

PROGRAM 9

RANDOM FOREST ALGORITHM

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
df = pd.read_csv("/content/drive/MyDrive/iris.csv")
```

```
df.head()
```

```
y = df[["species"]]
```

```
x = df.drop(["species"], axis=1)
```

```
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split  
(x, y, test_size=0.3, random_state=0)
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf = RandomForestClassifier(n_estimators=100)
```

```
clf.fit(x_train, y_train)
```

```
y_pred = clf.predict(x_test)
```

```
from sklearn.metrics import accuracy_score
```

```
score = accuracy_score(y_pred, y_test)
```

```
print(f"Accuracy : {score}%")
```

OUTPUT

Accuracy : 100%

AdaBoost (with default parameters)

```
from sklearn.ensemble import AdaBoostClassifier  
adb = AdaBoostClassifier()  
adb-model = adb.fit(x-train, y-train)  
y-pred = adb-model.predict(x-test)  
score = accuracy_score(y-pred, y-test)  
print(f"Accuracy: {score}")  
Accuracy: 0.977
```

AdaBoost (with Slighter parameters)

```
from sklearn.linear_model import LogisticRegression  
lrm = LogisticRegression()
```

```
adbhp = AdaBoostClassifier(n_estimators=150,  
                           estimator=lrm,  
                           learning_rate=0.1),  
model = adbhp.fit(x-train, y-train),  
y-pred = model.predict(x-test),  
score = accuracy_score(y-pred, y-test),  
print(f"Accuracy: {score}")
```

OUTPUT

Accuracy: 1.0

Adsf 23/5/24

LAB 10

KMEANS

```

import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import pandas as pd
import numpy as np

iris = datasets.load_iris()
X = pd.DataFrame(iris.data)
X.columns = ['Sepal-Length', 'Sepal-Width', 'Petal-Length',
             'Petal-Width']

y = pd.DataFrame(iris.target)
y.columns = ['Targets']

model = KMeans(n_clusters=3)
model.fit(X)

plt.figure(figsize=(14, 14))

plt.subplot(2, 2, 1)
colormap = np.array(['red', 'lime', 'black'])
plt.scatter(X.Petal-Length, X.Petal-Width, c=colormap[y['Targets']],
            s=40)

plt.title('Real clusters')
plt.xlabel('Petal Length')
plt.ylabel('Petal width')

plt.subplot(2, 2, 2)
plt.scatter(X.Petal-Length, X.Petal-Width, c=model.labels_),
            s=40)

plt.title('K-Means Clustering')
plt.xlabel('Petal Length')
plt.ylabel('Petal width')

```

PRINCIPAL COMPONENT ANALYSIS

LAB-11

```

import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
%matplotlib inline

```

```

from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
cancer.keys()
print(cancer['DESCR'])

```

```

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(df)
scaled_data = scaler.transform(df)

```

```

from sklearn.decomposition import PCA
pca = PCA(n_components=2)
pca.fit(scaled_data)

```

PCA (copy=True, n_components=2, whiten=False)

~~scaled_data.shape~~

```

plt.figure(figsize=(8,6))
plt.scatter(x_pca[:, 0], x_pca[:, 1], c=cancer['
```

plt.xlabel('First principal component')
 plt.ylabel('Second principal component')
 cmap='plasma')

Ans 30/5 fm